

# CS 211 project report: Remote access to OpenWRT

Zhehao Wang  
Department of Computer  
Science, UCLA  
zhehao@cs.ucla.edu

Haitao Zhang  
Department of Computer  
Science, UCLA  
haitao@cs.ucla.edu

Jeffrey Chen  
Department of Computer  
Science, UCLA  
jc4556@g.ucla.edu

## ABSTRACT

content...

## 1. GUIDELINE

abstract; introduction; background; design; implementation; evaluation/demo; discussions; related work; conclusion and future work.

## 2. INTRODUCTION

OpenWRT is a free, open-source, Linux-kernel based operating system (OS) for network-routing embedded systems. This OS is notable for being able to run on various types of devices and for simplifying cross-platform building of OpenWRT software packages, including fixes for devices no longer supported by the devices' manufacturers. OpenWRT receives regular updates, and allows basic router configuration, and installation of features through a package repository.

OpenWRT can be configured using a SSH command-line interface or a pre-packaged LuCI web interface. The SSH command-line interface is more suitable to professional and developer users, while a web interface is more friendly to common users, providing access to basic OpenWRT functions. However, SSH is tedious on a mobile device, because the user must perform all configurations via text modifications, and data and results cannot be interpreted in visual graphics. On the other hand, OpenWRT's existing remote access web interface is made for desktop web browsers, and not smart devices. On mobile smart devices, the in-browser interface is not scaled to the dimensions of the device's screen nor is it touch-friendly, hampering user comprehension and control of the interface. Additionally, depending on the device, loss of connection or suspension of the browser to another application can force renewal of the session or prevent retrieval of network information, causing issues for presenting real-time data and visualizations.

A native smart device application is more distribution-friendly for smart devices and can be designed on smart devices to be more user-friendly and streamlined to OpenWRT uses. Therefore, this project seeks to create a lightweight, easy-to-use generic Android application avail-

able for products running OpenWRT.

## 3. BACKGROUND

### 3.1 What is OpenWRT

OpenWRT [?, ?] is a Busybox/Linux based embedded platform which is developed following GPL license. It minimizes its own functions so that it fits for lots of memory constrained devices. Specifically, it builds the appropriate toolchain for devices, compiles appropriate kernel with patched and options, and provides software as IPKG packages.

### 3.2 OpenWRT System Structure

OpenWRT System Structure covers four aspects: directory structure, packages and external repositories, toolchain, and software architecture.

There are four key directories in the base: tools, toolchain, package and target. Tools and toolchain refer to common tools which will be used to build the firmware image, the compiler, and the C library.

In OpenWRT, almost all the packages are .ipk files. Users can choose what packages to install and what packages to uninstall based on their specific needs. Packages are either part of the main trunk or maintained out of the main trunk. For the second case, packages can be maintained by the package feeds system.

To compile the program for a particular architecture, the OpenWRT system will automatically create the toolchain during the cross-compilation process, simplifying the development tasks. However, if the toolchain needs to be created manually, OpenWRT also provides an easy way to configure the arguments.

Figure 1 shows the software stack of OpenWRT. We can see that the common embedded Linux tools such as uClibc, busybox, shell interpreter are used by OpenWRT.

### 3.3 How to Develop With OpenWRT

It is easy to port software to OpenWRT. Various fetching methods such as GIT, Subversion, CVS, HTTP, local source can be used to download package source. In a typical package directory, there should always be a

UCI	IPKG	User programs
Busybox		
uClibc		
Linux kernel		

**Figure 1: the software stack of OpenWRT**

package/<name>/Makefile. After running “make menuconfig”, the new package will show up in the menu; and after running “make”, the new package will be built.

### 3.4 Web-based Access to OpenWRT

In the OpenWRT system, some important features are provided: a built-in web server with CGI support, an SSH server and a package management tool. We can make use of these existing tools to build our Android application.

Apart from the basic tools, OpenWRT also supports LuCI [?], which is a browser-based tool to remotely configure the OpenWRT system. Specifically, it provides status visualization functions, system administration functions and network configurations. We would provide similar functions in our Android application.

Another useful reference is the Netgear [?] mobile application, which designs nice UIs that we can borrow ideas from. We need also to design the web-based access application UIs based on their work.

## 4. DESIGN

### 4.1 Application specific traffic statistics

One of the goals in this project was to identify and provide statistics for traffic belonging to specific user applications, such as quantifying incoming and outgoing traffic related to specific smartphone application and services, and determining the identities of the hosts generating the traffic going through the OpenWRT box. To achieve this goal, we considered the two following design approaches:

(A). Pre-install a service on each user device that would monitor application network usage information from the operating system and report the monitored statistics to a service on the OpenWRT box. This approach can likely generate very accurate results and provide each specific application’s statistical network usage. However, the major drawback to this method of collecting accurate traffic information is that it requires that each user voluntarily install and run this service on their devices, a requirement non-ideal for our goal: this service should not require additional actions and

permissions from user devices, and the OpenWRT box should be able to collect statistics with or without end user cooperation.

(B). Capture the traffic on the OpenWRT box, and run per-packet analysis. While less accurate than the first option, it is less intrusive to the user. The following information components from packet headers or payloads could be used to capture traffic:

- Source and destination IP addresses.

Popular service providers, such as Google and Facebook, own large chunks of IP addresses, and certain destination IP address ranges may correspond to servers for a specific application. By collecting such information and building a mapping from destination IP address to an application’s backend services, it is possible to infer which user application generated the traffic. The problem with this approach is that this mapping takes time to build, and dynamically changing the mapping in response to newly collected information is difficult, especially considering the fact that many popular services are using CDNs.

- Source and destination port number.

Certain services can be identified by IANA’s port number allocation [?], and this information is easily retrievable from the Internet. The problem with this approach is that the number of applications that can be identified purely by port number is very limited, as some applications will share the same ports. For example, the Youtube and Facebook apps on Android both go through port 443 (TLS).

- Application layer payload.

Decoding application payload and trying to find characteristic plain text is another way to identify which application a packet belongs to. This approach is promising, except it works only if such characteristic texts can be located and the application payload has not encrypted, which is not the case for most popular Android applications, such as Youtube or Facebook. Packet capturing of these Android applications shows that their traffic goes on top of TLS, and packet capturing of the activity of watching Youtube videos on a desktop indicates that the traffic goes on top of QUIC, which has encryption over UDP.

- Source and destination host names.

DNS host names often give information about the service provider, even when the service provider’s using CDNs. By doing a DNS reverse lookup on the IP addresses, we can find out the service’s domain name, thus infer what application’s generating the traffic. The assumption behind this approach is that most servers or CDN boxes have

a DNS name that corresponds to the application backend that they run. Our initial experiments suggest that this assumption is indeed true for popular applications like Youtube and Facebook, though mapping a DNS domain to a certain service may not be straightforward. For example, Facebook app on Android talks to both *\*.facebook.com* domain, and *\*.fbcdn.net* (owned by the CDN provider, Akamai Technologies) domain, and Youtube app talks to both *\*.google.com*, and *\*.1e100.net* to another Google owned domain.

Given this caveat, the authors chose the last approach since it is most applicable than other options listed above. For the proof-of-concept implementation, the mapping from domain names to service names is statically configured. It is also worth mentioning that this approach will limit our statistics to per service provider, rather than per exact application (for example, the traffic application is unable to differentiate Google Hangout application traffic from Youtube application traffic), but the authors believe this approach is enough for the purpose of providing statistics on the router end that provides a network administrator a reasonable amount of information to determine which applications on which devices are generating how much traffic.

## 4.2 Mobile Application

This sub-section describes the design of the OpenWRT remote access mobile application. Since this is a proof of concept, we only designed and implemented the application on Android platform.

### 4.2.1 Architecture

Figure 2 shows the architecture of the Android OpenWRT remote access mobile application. It mainly contains four sub modules: activity module, network communication module, graphical user interface module and response result parser module. Among these four modules, activities combine the other three modules together to present result to users and take users' inputs.

### 4.2.2 Network Communication

Since the application needs to communicate with OpenWRT backend server, it's necessary to provide a network communication module to handle all the network operation, so that other modules can simply use its APIs without worrying about the details.

LuCI backend and application specific traffic analysis backend both provide http-based interfaces. To communicate with these two backends, the network communication module needs to be able to send "http POST" and "http GET" requests and receives responses. A key part is to build proper URLs according to backends' requirements. The only difference between "http POST" and "http GET" is that "http POST" carries some parameters in the URL. For example, a "http

POST" request to communicate with LuCI backend is "http://<IP address of the OpenWrt Box>:<server port>/cgi-bin/luci/;stok=<stok id>"; a "http GET" request to communicate with application specific traffic analysis backend is "http://<IP address of the OpenWrt Box>:<server port>/output.txt"

### 4.2.3 Graphical User Interface

The goal of the graphical user interface design was to simplify the user interface on a smartphone. To that end, the limitations of the LuCI web interface were studied to provide design guidelines. The first issue analyzed was that LuCI had an issue in its navigation on smartphones. To navigate through the application categories, the user needed to select a category in the navigation menu, then select a subcategory from the dropdown menu. Additionally, changing subcategories within the same category still required selecting the overarching category again. Therefore, a design goal of the application would be to maintain the current category and simply swap subcategories.

Another LuCI WebView issue was that unsaved changes would be tracked in a session until committed. Tracking unsaved changes on a web browser can result in session complications, depending on browser settings for caching. Furthermore, the WebView relied on in-browser scripting to provide functional elements, which is a dependency that can be optimized. Therefore another aspect of our design was to make all actions atomic and contained to the screen they are accessed on, to avoid carrying changes. To make all actions atomic, all functional elements in the original WebView would be rebuilt natively in Android.

Based on the LuCI framework, the designed Android application's user interface screens consisted of a login screen, then three major categories: status, network, and system. Each category then presented a subnavigation menu that persisted until another major category was selected, allowing users to move more freely within same category.

Each separate screen in the subcategories of the major categories was designed to maintain discrete actions and information. Rather than having multiple configuration forms and submission buttons in the same screen, the screen would be limited to at most one form each, with other forms being accessible on a new screen that is linked to by a list on the current screen.

## 5. IMPLEMENTATION

### 5.1 Application specific traffic statistics implementation

Based on the design described in section 4.1, the proof of concept implementation pipelines the output of *tcpdump* to a custom backend. The backend then reads the source and destination IP addresses, tries to use *nslookup* to find domain names by DNS reverse




Figure 2: Android Application Architecture

lookup, and matches the domain names with statically configured characteristic strings. If a match is found, the number of bytes is added for corresponding service provider name and corresponding host; otherwise, the traffic is considered identified. The backend writes statistics in JSON format to a file, and the file is served by a simple web server. In the statistics, incoming and outgoing traffic, as well as different IP addresses within the local subnet, is differentiated. An example of the JSON object is given below, in which the numbers are number of bytes per IP address and application. We use Python for the simple backend and web server, for easy prototyping purposes.

```

1      {"outgoing":
2      {"google": {"10.0.5.15": 812},
3      "facebook": {"10.0.5.15": 622}},
4      "incoming":
5      {"google": {"10.0.5.15": 148865},
6      "facebook": {"10.0.5.15": 76947}}
7      }

```

Basic optimizations for improving the performance of the backend include introducing the following:

- A cache for IP and domain mapping, so that instead of calling *nslookup* each time a packet arrives, *nslookup* is tried only once for each IP address within a given time.
- A minimum waiting interval of one minute between writing statistics back, so that the number of file writes is limited.

## 5.2 Android Application Implementation

### 5.2.1 Graphical User Interface

The development of the user interface considered native Android solutions to the navigation. Typically, Android activities are single interface screens. Inside the Android activities, Android fragments can be used

to provide different tab screens. Therefore, tab navigation was implemented with Android activities for the categories, and Android fragments for the subcategories and forms.

However, the implementation of the user interface reached complications. The major problem leading to complications was that in using the LuCI backend, content extraction using HTML retrieved whole WebView pages. JSON would extract smaller elements, but required specific queries, and would make code reuse difficult. Using both required parsing, and the extent of parsing needed for each category and subcategory complicated the implementation, taking more development time than available. The issue of parsing elements also complicated the implementation of buttons and other interactive objects, which needed to be re-made natively. Therefore, the graphical user interface was not completed in time.

## 6. EVALUATION

### 6.1 Application specific traffic statistics backend evaluation

This section describes the disk, CPU and memory usage of the Python backend for application specific traffic analysis.

Hard drive volume's usually quite limited for an OpenWRT box. The backend analyzer has dependency on *opkg* packages *tcpdump*, *python-light* and *python-codecs*; and the associated http server has dependency on *python-logging* and *python-ssl* (optional, the module would not give out warnings if installed). On our test TP-Link WDR4300 router flashed with OpenWRT 15.05, 500KB disk space still remains for mount point `/`, after all the necessary dependencies are installed. The mount point `/tmp` has over 60MB temporary space, and we store our *tcpdump* file, and statistics output file under mount point `/tmp`.

CPU and memory usage on the OpenWRT boxes's measured with users behind the AP watching a Youtube video. When only one user's behind the AP watch-

ing a Youtube video at 360p, our Python backend uses 8% 22% CPU, tcpdump uses 1% 10% depending on whether the video is pre-buffered long enough. Both tcpdump and Python backend use 5% of the virtual memory. With two users behind the AP, the Python backend CPU usage grows by 7% on average, and we did not see tcpdump CPU usage ramp up. Since the communication between tcpdump and the Python backend is through file system interface, IO is more likely the bottleneck here than CPU.

## 6.2 Demo description

A screen recording of the current Android application is available at <http://memoria.ndn.ucla.edu/openwrt2.mp4>. This demo uses a TP-Link WDR4300 router with LuCI and the application specific analysis backend installed. In the demo, the phone application first configures the SSID of the router from “OpenWRT” to “OpenWRT123”, and reconnects to WiFi using the changed SSID. Then the phone generates Youtube traffic by visiting its webpage, and demonstrates that the incoming traffic labeled “Google” grows by 4MB correspondingly. While the previous demo uses WebView to browse and configure the wireless AP, another screen recording at <http://memoria.ndn.ucla.edu/openwrt1.mp4> demonstrates our initial attempts of moving towards JSON and html parsing, using native Java code instead of executing JS in the WebView, and using Android UI widgets instead of relying on the WebView CSS for styling.

## 7. DISCUSSIONS

## 8. FUTURE WORK

The future work for this project is identified as the following:

- Application specific traffic analysis module: rewriting the backend in C, remove dependency on tcpdump and communication via file system, differentiating traffic inside and outside the local subnet, creating traffic statistics based on time, and integration with LuCI backend.
- Android UI: migrating to using native Java code and native Android UI widgets instead of the WebView.
- Backend pattern of communication simplification: instead of passing back html scripts from backend and parse the html script in the application, the backend should just reply with JSON objects including the Android application’s queried objects.

## 9. CONCLUSION

In this project we designed and implemented an Android application for configuring an OpenWRT AP, and a Python backend on the AP to provide per-user, per-application statistics.

The Android application communicates with LuCI via http requests, and calls its functions to authenticate a user and configure the AP. The per-application statistics backend gathers input from *tcpdump*, does DNS reverse lookup to decide the DNS names for packet’s source and destination IP addresses, and use the DNS name to decide which service provider’s application generated the traffic. The Android application also talks with the per-application statistics backend, and provides a visualization of its statistics.

We provided demonstrations of the initial frontend application, the UI designs for future versions of the application, and basic profiling results of the per-application statistics backend.

## 10. TIMELINE

A rough timeline for the project is given in table 1.

**Table 1: Project timeline**

Week No.	Task
5, 6 (first half)	Implement status/statistics visualization module
6 (second half), 7	Implement network configuration module
8, 9 (first half)	Implement system administration module
9 (second half), 10	Prepare final report and presentation